

NetFPGA Summer Course



Presented by:

**Andrew W Moore, Noa Zilberman, Gianni Antichi
Stephen Ibanez, Marcin Wojcik, Jong Hun Han,
Salvator Galea, Murali Ramanujam, Jingyun Zhang,
Yuta Tokusashi**

**University of Cambridge
July 24 – July 28, 2017**

<http://NetFPGA.org>

Previously Covered

- The NetFPGA platform
- Network Review
- The Base Reference Router
- The Life of a Packet Through the NetFPGA
- Infrastructure
- Examples of Using NetFPGA
- Example Project: Crypto Switch
- Simulation and Debug

Tutorial Outline

- **Register Infrastructure**
 - Explain register system
 - Use AXI Lite registers modules to implement register
 - Add register access stimulus to define Crypto Switch encryption key
 - Interface with software
 - Hardware tests
- **Build and Test Hardware**
 - Build
 - Explanation of Hardware Tests
 - Write Hardware Tests
 - Program the board

Section I: Register Infrastructure

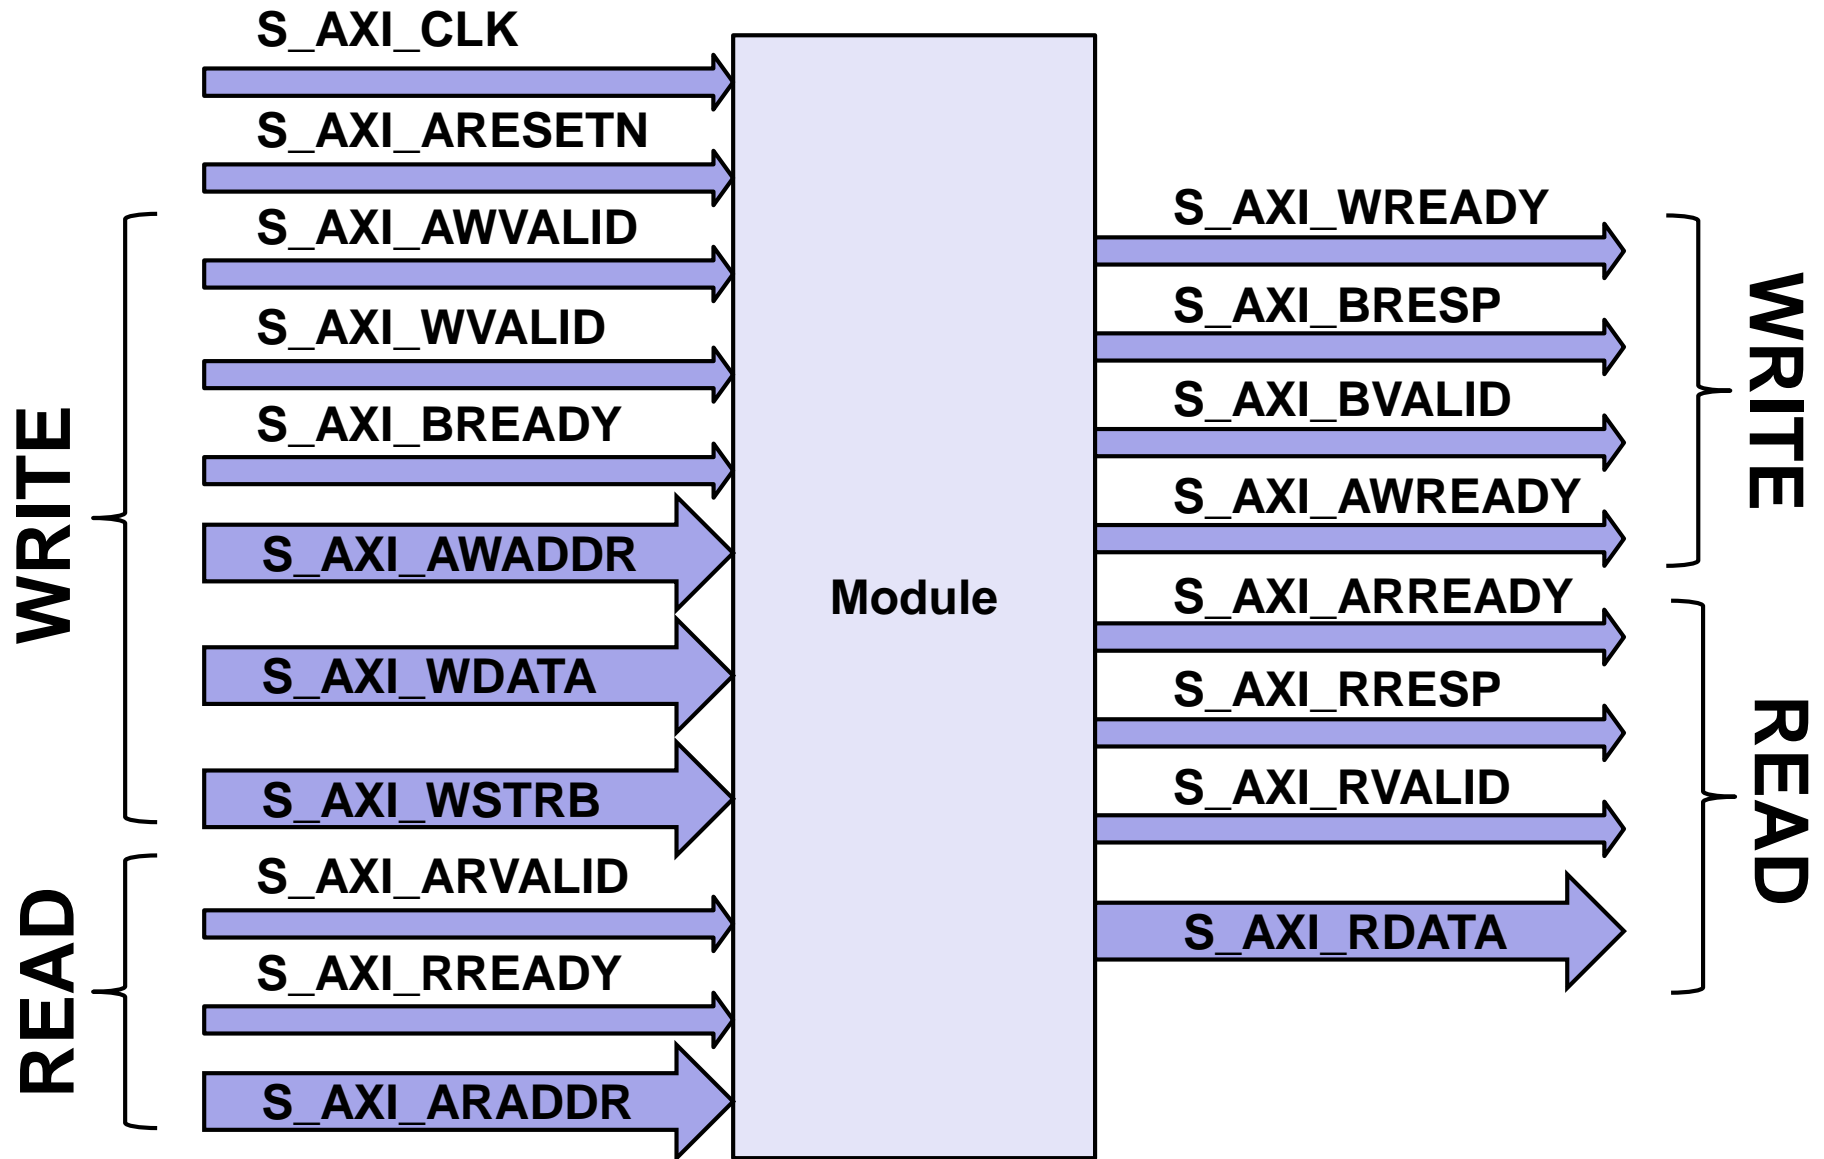
Specifying the Key via a Register

- **Set the key via a register**
 - Instead of a constant value
- **Requires understanding the registers system 😊**
- **Registers system:**
 - Automatically generated
 - Implementing registers in a module
 - Use automatically generated `cpu_regs` module
 - Need to implement the registers' functional logic

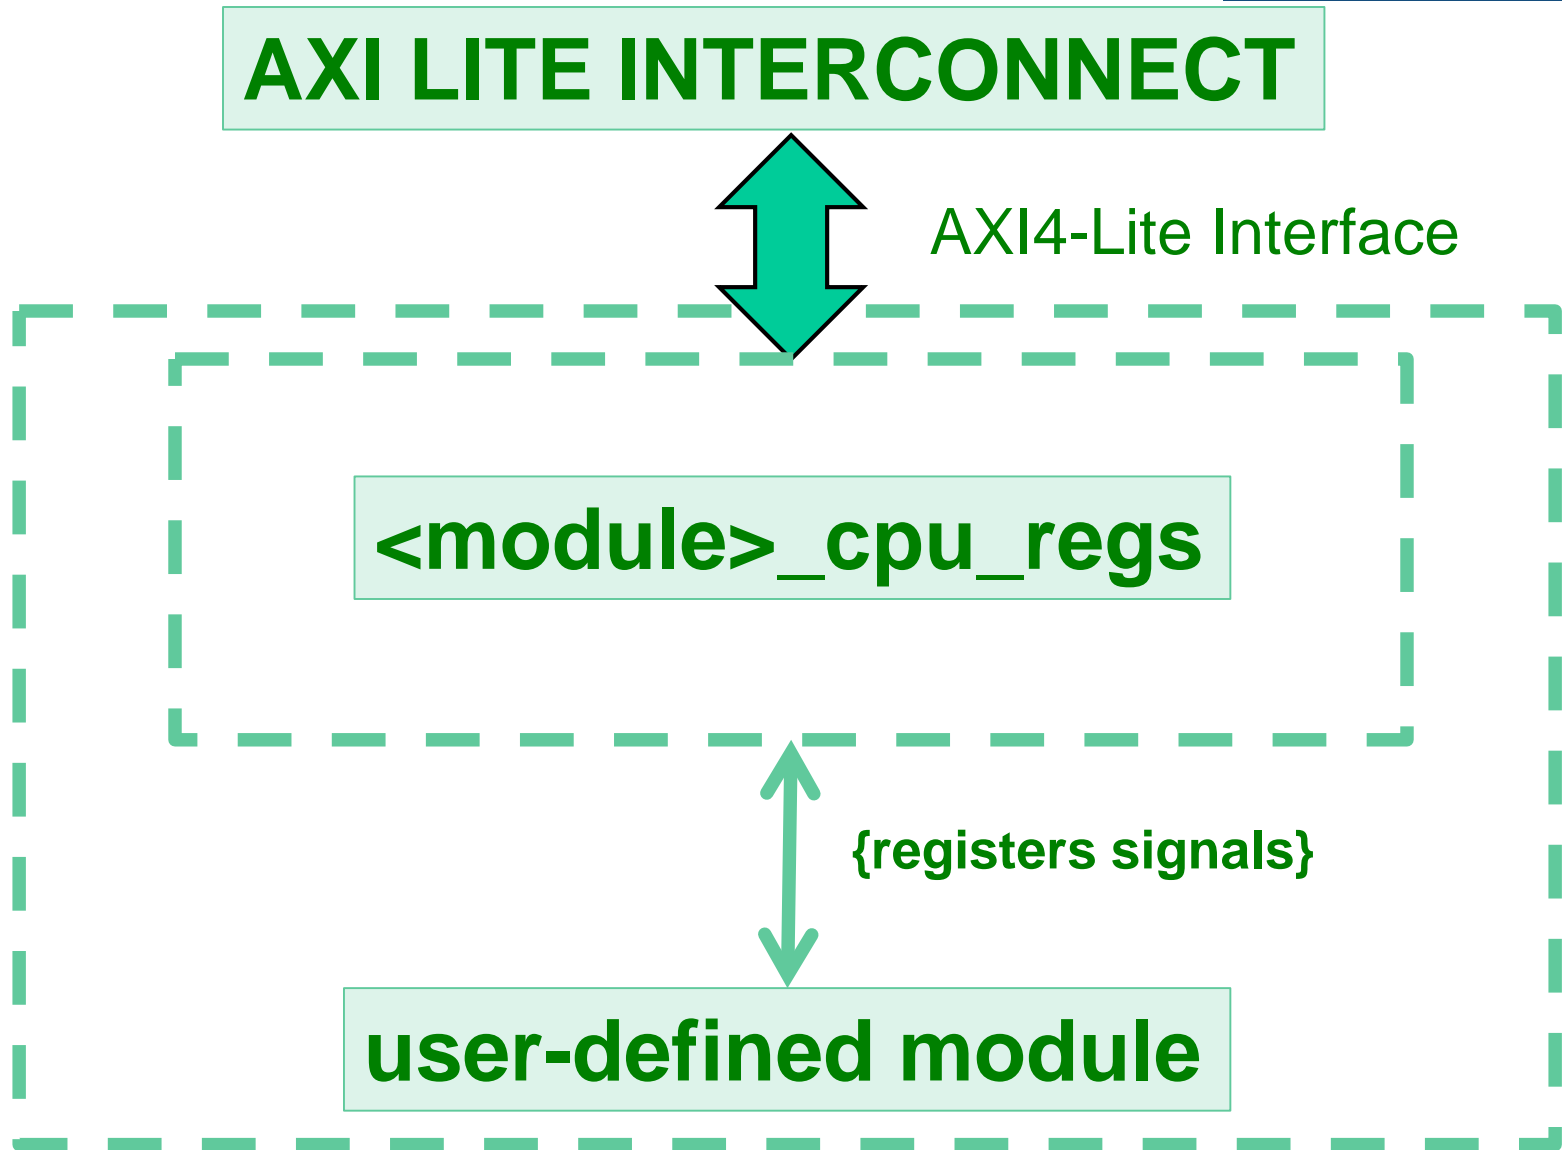
Registers bus

- **Yesterday we learnt that packets stream follows the AXI4-Stream paradigm**
- **Register communication follows the AXI4-Lite paradigm**
- **The AXI4-Lite interface provides a point-to-point bidirectional interface between a user Intellectual Property (IP) core and the AXI Interconnect**

Register bus (AXI4-Lite interface)



Register bus



Registers – Module generation

- Spreadsheet based (xls / csv)
- Defines all the registers you intend to support and their properties
- Generates a python script (regs_gen.py), which generates the outputs
- Location: \$SUME_FOLDER/tools/infrastructure

	Generate Registers			OS:	Windows					
Block	Register Name	Address	Description	Type	Bits	Endian Type	Access Mode	Valid for sub-modules	Default	Constraints, Remarks
IP_name	Init	NA	When triggered, the module will perform SW reset	Global	0	Little		sub_ip_name		
IP_name	ID	0	The ID of the module, to make sure that one accesses the right module	Reg	31:0	Little	RO	sub_ip_name	32'h0000DA03	
IP_name	Version	4	Version of the module	Reg	31:0	Little	RO	sub_ip_name	32'h1	
IP_name	Flip	8	The register returns the opposite value of what was written to it	Reg	31:0	Little	RWA	sub_ip_name	32'h0	Returned value is at reset 32'hFFFFFF
IP_name	CounterIn	C	Incoming Packets Counter	Reg	31:0	Little	ROC	sub_ip_name	32'h0	
	CounterIn		Number of Incoming packets through the	Field	30:0		ROC	opl	31'h0	
	CounterInOvf		Counter Overflow indication	Field	31		ROC	opl	1'b0	
IP_name	CounterOut	10	Outgoing Outgoing Packets Counter	Reg	31:0	Little	ROC	sub_ip_name	32'h0	
	CounterOut		Number of Outgoing packets through the	Field	30:0		ROC	opl	31'h0	
	CounterOutOvf		Counter Overflow indication	Field	31		ROC	opl	1'b0	
IP_name	Debug	14	Debug Register, for simulation and debug	Reg	31:0	Little	RWA	sub_ip_name	32'h0	
IP_name	EndianEg	18	Example big endian register	Reg	31:0	Big	RWA	sub_ip_name	32'h0	

Registers – Module generation

	Generate Registers			OS:	Windows					
Block	Register Name	Address	Description	Type	Bits	Endian Type	Access Mode	Valid for sub-modules	Default	Constraints, Remarks
IP_name	Init	NA	When triggered, the module will perform SW reset	Global	0	Little		sub_ip_name		
IP_name	ID	0	The ID of the module, to make sure that one accesses the right module	Reg	31:0	Little	RO	sub_ip_name	32'h0000DA03	
IP_name	Version	4	Version of the module	Reg	31:0	Little	RO	sub_ip_name	32'h1	
IP_name	Flip	8	The register returns the opposite value of what was written to it	Reg	31:0	Little	RWA	sub_ip_name	32'h0	Returned value is at reset 32'hFFFFFF
IP_name	CounterIn	C	Incoming Packets Counter	Reg	31:0	Little	ROC	sub_ip_name	32'h0	
	CounterIn		Number of Incoming packets through the	Field	30:0		ROC	opl	31'h0	
	CounterInOvf		Counter Overflow indication	Field	31		ROC	opl	1'b0	
IP_name	CounterOut	10	Outgoing Outgoing Packets Counter	Reg	31:0	Little	ROC	sub_ip_name	32'h0	
	CounterOut		Number of Outgoing packets through the	Field	30:0		ROC	opl	31'h0	
	CounterOutOvf		Counter Overflow indication	Field	31		ROC	opl	1'b0	
IP_name	Debug	14	Debug Register, for simulation and debug	Reg	31:0	Little	RWA	sub_ip_name	32'h0	
IP_name	EndianEg	18	Example big endian register	Reg	31:0	Big	RWA	sub_ip_name	32'h0	

Registers – Module generation

Access Modes:

- **RO - Read Only (by SW)**
- **ROC - Read Only Clear (by SW)**
- **WO - Write Only (by SW)**
- **WOE - Write Only Event (by SW)**
- **RWS - Read/Write by SW**
- **RWA - Read/Write by HW and SW**
- **RWCR - Read/Write clear on read (by SW)**
- **RWCW - Read/Write clear on write (by SW)**

Registers – Module generation

Endian Mode:

- **Little Endian – Most significant byte is stored at the highest address**
 - Mostly used by CPUs
- **Big Endian - Most significant byte is stored at the lowest address**
 - Mostly used in networking
 - e.g. IPv4 address

Registers – Generated Modules

- **<module>_cpu_regs.v – Interfaces AXI-Lite to dedicated registers signals**
To be placed under under <core name>/hdl
- **<module>_cpu_regs_defines.v – Defines per register: width, address offset, default value**
To be placed under under <core name>/hdl
- **<module>_cpu_template.v – Includes template code to be included in the top core Verilog.**
This file can be discarded after updating the top core verilog file.

Registers – Generated Modules

Same contents as `<module>_cpu_regs_defines.v`, but in different formats, used by software, build and test harness:

- `<module>_regs_defines.h`
To be placed under `<core name>/data`
- `<module>_regs_defines.tcl`
To be placed under `<core name>/data`
- `<module>_regs_defines.txt` – used by test harness
To be placed under `<core name>/data`

Adding Registers Logic - Example

- Usage examples:

```
always @(posedge axi_aclk)
  if (~resetn_sync) begin
    id_reg <= #1    `REG_ID_DEFAULT;
    ip2cpu_flip_reg <= #1    `REG_FLIP_DEFAULT;
    pktin_reg <= #1    `REG_PKTIN_DEFAULT;
  end
  else begin
    id_reg <= #1    `REG_ID_DEFAULT;
    ip2cpu_flip_reg <= #1    ~cpu2ip_flip_reg;
    pktin_reg <= #1    pktin_reg_clear ? 'h0 :
                        pkt_in ? pktin_reg + 1: pktin_reg ;
  end
end
```

NetFPGA-Host Interaction

- Register reads/writes via ioctl system call
- Useful command line utilities

```
cd $APPS_FOLDER/sume_riffa_v1_0_0/  
./rwaxi -a 0x44010000  
./rwaxi -a 0x44010000 -w 0x1234
```

You must program the FPGA and load the driver before using these commands!

**Can I collect the registers addresses in
a unique .h file?**

NetFPGA-Host Interaction

- Need to create the `sume_register_defines.h` file
 - `cd $NF_DESIGN_DIR/hw`
 - `make reg`
- The `sume_register_defines.h` file will be placed under `$NF_DESIGN_DIR/sw/embedded/src`

NetFPGA-Host Interaction

Required steps:

- Generate .h file per core
 - Automatically generated by the python script
- Edit \$NF_DESIGN_DIR/hw/tcl/
\$NF_PROJECT_NAME_defines.tcl
 - Indicate the address mapping you use
- Edit \$NF_DESIGN_DIR/hw/tcl/
export_registers.tcl
 - Indicate the location of all IP cores used
 - Default path assumed is under \lib\hw\cores

NetFPGA-Host Interaction

- `sume_register_defines.h` is automatically generated when creating a project
 - Using NetFPGA TCL scripts, the `.h` file will match the hardware
 - Note that changes in the GUI will not be reflected!
- Post implementation, for the SDK, use `$NF_DESIGN_DIR/hw/tcl/export_hardware.tcl`
 - Uses vivado's export
 - Does not include the registers list, only memory map


Testing Registers with Simulation

Testing Registers with Simulation

- **nftest_regwrite(address, value)**
 - nftest_regwrite(0x44010008, 0xABCD)
- **nftest_regread(address)**
 - nftest_regread(0x44010000)
- **nftest_regread_expect(address, expected_value)**
 - nftest_regread_expect(0x44010000, 0xDA01)
- **Can use registers names**
 - nftest_regread(SUME_INPUT_ARBITER_0_ID)
- **Use within run.py**
- **You don't need to edit any other file**


Simulating Register Access

1. Define register stimulus
2. The testbench executes the stimulus



3. Simulation accesses are written to a log file

4. A script can compare expected and actual values
And declare success or failure



Legend:

- DUT: Design Under Test
- stim: stimulus
- tb: testbench
- sim: simulation

Registers Stimulus (1)

```
cd $NF_DESIGN_DIR/test/  
less reg_stim.axi
```

- An example of write format :

Ten DWORD writes to nic_output_port_loopup interface. Each waits for completion.

77000000	deadcode	f	-.
77000004	acce55ed	f	-.
77000008	add1c7ed	f	-.
7700000c	ca0ebabe	f	-.
77000010	c0dedead	f	-.
77000014	55edacce	f	-.
77000018	babeca1e	f	-.
7700001c	abcde9ab	f	-.
77000020	cde2abcd	f	-.
77000024	e4abcde3	f	-.

Address

Data

Byte Enable strobe

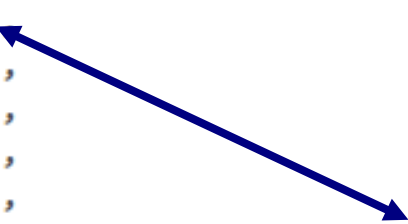
with other useful information like, time, barriers etc..

Registers Stimulus (2)

```
cd $NF_DESIGN_DIR/test/both_testreg_crypto  
less reg_stim.axi
```

- An example read format :

```
# Ten DWORD quick reads from the nic_output_port_loopup interface (without waits.)  
-, -, -, 77000000  
-, -, -, 77000004,  
-, -, -, 77000008,  
-, -, -, 7700000c,  
-, -, -, 77000010,  
-, -, -, 77000014,  
-, -, -, 77000018,  
-, -, -, 7700001c,  
-, -, -, 77000020,  
-, -, -, 77000024. # Never wrap addresses until after WAIT flag!
```



Address

with other useful information like, time, barriers etc..

Registers Access Log

```
cd $NF_DESIGN_DIR/test/both_testreg_crypto  
less reg_stim.log
```

WRITE

77000000	<-	DEADC0DE	(OKAY)	# 1335 ns
77000004	<-	ACCE55ED	(OKAY)	# 1405 ns
77000008	<-	ADD1C7ED	(OKAY)	# 1475 ns
7700000C	<-	CA0EBABE	(OKAY)	# 1545 ns
77000010	<-	C0DEDEAD	(OKAY)	# 1615 ns
77000014	<-	55EDACCE	(OKAY)	# 1685 ns
77000018	<-	BABECA1E	(OKAY)	# 1755 ns
7700001C	<-	ABCDE9AB	(OKAY)	# 1825 ns
77000020	<-	CDE2ABCD	(OKAY)	# 1895 ns
77000024	<-	E4ABCDE3	(OKAY)	# 1965 ns
77000000	<-	DEADC0DE	(OKAY)	# 2035 ns
77000004	->	ACCE55ED	(OKAY)	# 2095 ns
77000008	->	ADD1C7ED	(OKAY)	# 2155 ns
7700000C	->	CA0EBABE	(OKAY)	# 2215 ns
77000010	->	C0DEDEAD	(OKAY)	# 2275 ns
77000014	->	55EDACCE	(OKAY)	# 2335 ns
77000018	->	BABECA1E	(OKAY)	# 2395 ns
7700001C	->	ABCDE9AB	(OKAY)	# 2455 ns
77000020	->	CDE2ABCD	(OKAY)	# 2515 ns
77000024	->	E4ABCDE3	(OKAY)	# 2575 ns

READ

Time

Replacing Static Key

- **In the crypto project, replace the static key with the key from the registers**
 - Use a RWA or a RWS register for your key
 - Assign a default value through the RW default register
- **Hint: Is this a Big or a Little Endian register?**
- **Check in simulation if the system still works correctly**
 - Use both_testreg_crypto

Section II: Build and Test Hardware

Synthesis

- **To synthesize your project:**

```
cd $NF_DESIGN_DIR  
make
```

Hardware Tests

- **Test compiled hardware**
- **Test infrastructure provided to**
 - Send Packets
 - Check Counters
 - Read/Write registers
 - Read/Write tables

Python Libraries

- **Start packet capture on interfaces**
- **Clear all tables in hardware**
- **Create packets**
 - MAC header
 - IP header
 - PDU
- **Read/Write registers**
- **Read/Write reference router tables**
 - Longest Prefix Match
 - ARP
 - Destination IP Filter
- **The same libraries used in the simulation infrastructure...**

Creating a Hardware Test

Useful functions:

Packet generation:

`make_IP_pkt(...)` – see [wiki](#)

`encrypt_pkt(key, pkt)`

`decrypt_pkt(key, pkt)`

Packet transmission/reception:

`nftest_send_phy(interface, pkt)`

`nftest_expect_phy(interface, pkt)`

`nftest_send_dma(interface, pkt)`

`nftest_expect_dma(interface, pkt)`

Register access:

`nftest_regwrite(addr, value)`

`nftest_regread_expect(addr, expect)`

Understanding Hardware Test

- `cd $NF_DESIGN_DIR/test/both_crypto_test`
- `vim run.py`
- “isHW” indicates HW test
- “connections/conn” file declares the physical connections

```
nf0:eth1
nf1:eth2
nf2:
nf3:
```
- “global/setup” file defines the interfaces

```
proc = Popen(["ifconfig","eth2","192.168.101.1"],
stdout=PIPE)
```

Your task:

- Remember to source the settings.sh file
- Edit run.py to create your test
- Edit setup and conn files

Running Hardware Tests

- **Use command `nf_test.py`**


- Required Parameter

- `sim hw` or `both` (right now only use `hw`)

- Optional parameters

- `--major <major_name>`
 - `--minor <minor_name>`

`both_crypto_test`



- **Run the command**

`nf_test.py hw --major testreg --minor crypto`

Running Hardware Tests

- **Having problems?**
- **Take advantage of the wiki!**
<https://github.com/NetFPGA/NetFPGA-SUMEPublic/wiki/Hardware-Tests>
 - Detailed explanations
 - Tips for debug

Recap

Build a complete NetFPGA design

Learn:

- **Module creation (Verilog)**
- **Reference pipeline integration**
- **Verification via simulation**
- **Verification via hardware tests**
- **Interaction with software**

...and now let's program the board!!!

Program the NetFPGA

Several options:

- **Program the bit file using Vivado's Hardware Manager**
- **Load a bit file for FPGA programming using impact script**
`~/NetFPGA-SUME-live/tools/scripts/load_bitfile.py \`
`-i $DESIGN_DIR/bitfiles/crypto_switch.bit`
- **Use Xilinx Microprocessor Debugger (XMD)**
`xmd`
`fpga -f <filename.bit>`

WHILE YOU WAIT.... Here is one we built earlier:

```
~/NetFPGA-SUME-live/tools/scripts/load_bitfile.py -i \  
~/NetFPGA-SUME-live/projects/crypto_switch_solution_day2/bitfiles/  
crypto_switch.bit
```

Loading the driver

- **Compile SUME driver:**
 - `cd $DRIVER_FOLDER`
 - `make`
 - `make install`
 - `insmod sume_riffa.ko`
- **Must reset the computer after programming the FPGA**
 - For proper detection and enumeration of PCIe
- **If you already had a running board**
 - `cd $SUME_FOLDER/tools/scripts/reconfigure`
 - `source pci_rescan_run.sh`
 - rescans the pcie bus (does not always succeed)

All in one

- `cd $SUME_FOLDER/tools/scripts/`
- `./run_load_image <filename.bit>`

- **Programs the board**
- **Rescans PCIe**
- **Loads the driver**
- **Brings up the 10G interfaces**

Section III: Projects

Thoughts for Developers

- **Build Modular components**
 - Describe shared registers
 - Consider how modules would be used in larger systems
- **Define functionality clearly**
 - Through regression tests
 - With repeatable results
- **Disseminate projects**
 - Post open-source code
 - Document projects on Web, Wiki
- **Expand the community of developers**
 - Answer questions on the Email list
 - Collaborate with your peers to build new applications

Finding Your Dream Team

- **Define your goals: what is your research agenda? What are you trying to achieve?**
- **What is your assessment criteria?**
 - Functionality, performance, programmability, study curve etc.
- **Pick a suitable language**
 - Verilog, P4, C#,... - each fits a different criteria.
- **Find other people with the same interests**
 - Use the list.
- **Form a team!**

Project Ideas for the NetFPGA

- Drop 1-in- N packet module
- Lower-latency switch
- On-platform power measurements
- Event-triggered network monitor
- Adding INT support
- SRAM / DRAM based output queues
- Rate-limiting module
- Input / Output scheduler
- A simple OpenFlow switch
- 40G Port
-

Previous Camp Projects

- **Circuit breaker (HFT transactions monitor)**
- **Payload censoring block**
- **Simple cool packet generator**
- **Histogram of size of packets**
- **Improving the latency of OpenFlow switch**
- **ACL Firewall**
- **Stateless NAT**
- **Dynamic port-based firewall (SW project)**

Planning Ahead

- This afternoon is about P4→ NetFPGA
- This evening is time to try hw testing & synthesis of crypto_switch design if you didn't manage that

For tomorrow:

- Select your group, leader & project
- Specific, Realistic
- Time-bounded
- Modest is good.....
- Discuss with us!
- Presents your planned project to the class

Acknowledgments (I)

NetFPGA Team at University of Cambridge (Past and Present):

Andrew Moore, David Miller, Muhammad Shahbaz, Martin Zadnik, Matthew Grosvenor, Yury Audzevich, Neelakandan Manihatty-Bojan, Georgina Kalogeridou, Jong Hun Han, Noa Zilberman, Gianni Antichi, Charalampos Rotsos, Hwanju Kim, Marco Forconesi, Jinyun Zhang, Bjoern Zeeb, Robert Watson, Salvator Galea, Marcin Wojcik, Diana Andreea Popescu, Murali Ramanujam

NetFPGA Team at Stanford University (Past and Present):

Nick McKeown, Glen Gibb, Jad Naous, David Erickson, G. Adam Covington, John W. Lockwood, Jianying Luo, Brandon Heller, Paul Hartke, Neda Beheshti, Sara Bolouki, James Zeng, Jonathan Ellithorpe, Sachidanandan Sambandan, Eric Lo, Stephen Gabriel Ibanez

All Community members (including but not limited to):

Paul Rodman, Kumar Sanghvi, Wojciech A. Koszek, Yahsar Ganjali, Martin Labrecque, Jeff Shafer, Eric Keller, Tatsuya Yabe, Bilal Anwer, Yashar Ganjali, Martin Labrecque, Lisa Donatini, Sergio Lopez-Buedo, Andreas Fiessler, Robert Soule, Pietro Bressana, Yuta Tokusashi
Steve Wang, Erik Cengar, Michael Alexander, Sam Bobrowicz, Garrett Aufdemberg, Patrick Kane, Tom Weldon
Patrick Lysaght, Kees Vissers, Michaela Blott, Shep Siegel, Cathal McCabe

Acknowledgements (II)



UNIVERSITY OF
CAMBRIDGE

EPSRC

Pioneering research
and skills



The Leverhulme Trust



ALGO-LOGIC

